

MATHENDRA ENGINEERING COLLEGE FOR WOMEN

NAMESHVARVAS

CLASSIVEAR-CSE

SUBJBM(AI)

REGISTER NO:611419104019

#libraries

import pandas as pd

import numpy as npp

import matplotlib.pyplot as plt

%matplotlib inline

#load dataset

df = pd.read_csv(r"/content/Churn_Modelling.csv")

df.head(10)

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43
5	6	15574012	Chu	645	Spain	Male	44
6	7	15592531	Bartlett	822	France	Male	50
7	8	15656148	Obinna	376	Germany	Female	29
8	9	15792365	He	501	France	Male	44
9	10	15592389	H?	684	France	Male	27
	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember		
0	2	0.00	1	1	1		
1	1	83807.86	1	0	1		
2	8	159660.80	3	1	0		
3	1	0.00	2	0	0		
4	2	125510.82	1	1	1		
5	8	113755.78	2	1	0		
6	7	0.00	2	1	1		
7	4	115046.74	4	1	0		
8	4	142051.07	2	0	1		
9	2	134603.88	1	1	1		
	EstimatedSalary	Exited					
0	101348.88	1					
1	112542.58	0					
2	113931.57	1					

```
3      93826.63      0
4      79084.10      0
5      149756.71      1
6      1006280      0
7      119346.88      1
8      74940.50      0
9      71725.73      0
```

```
df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column              Non-Null Count  Dtype
---  -
0   RowNumber           10000 non-null  int64
1   CustomerId          10000 non-null  int64
2   Surname             10000 non-null  object
3   CreditScore         10000 non-null  int64
4   Geography           10000 non-null  object
5   Gender              10000 non-null  object
6   Age                10000 non-null  int64
7   Tenure              10000 non-null  int64
8   Balance             10000 non-null  float64
9   NumOfProducts      10000 non-null  int64
10  HasCrCard           10000 non-null  int64
11  IsActiveMember     10000 non-null  int64
12  EstimatedSalary    10000 non-null  float64
13  Exited              10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 11+ MB

# Visualizations
# Univariate Analysis

import seaborn as sns
sns.kdeplot(df['CreditScore'])

<matplotlib.axes._subplots.AxesSubplot at 0x7fc4a0cd2790>
```

```
#S1 - Variate Analysis
```


```
plt.bar(df.CustomerId, df.CreditScore)
```

```
plt.title('CreditScore')
```

```
plt.xlabel('CustomerId')
```

```
plt.ylabel('CreditScore')
```

```
Text(0, 0.5, 'CreditScore')
```



```
sns.lmplot(x='Tenure', y='Balance', data=df, hue='Exited', size=8)
/usr/local/lib/python3.7/dist-packages/seaborn/regression.py:581:
UserWarning: The `size` parameter has been renamed to `height`; please
update your code
warnings.warn(msg, UserWarning)
<seaborn.axisgrid.FacetGrid at 0x7fc4d49e2d0>
```




```
#Multi - Variate Analysis
```

```
ax =
```

```
df[["CreditScore","Age","Tenure","Balance"]].plot(figsize=(80,40))
```

```
ax.legend(loc='center left', bbox_to_anchor=(1, 0.5)):
```



```
df.isnull().sum()
RowNumber      0
CustomerId      0
Surname         0
CreditScore     0
Geography      0
Gender          0
Age            0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard      0
IsActiveMember  0
EstimatedSalary 0
Exited         0
dtype: int64
plt.figure(figsize=(15,3))
sns.heatmap(df.corr(),annot=True,cmap='BuPu')
plt.show()
```

$$df.head()$$

```
1      0      1      112542.58      0
2      1      0      113931.57      1
3      0      0      93826.63      0
4      1      1      79084.10      0
```

```
df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column              Non-Null Count  Dtype
---  -
0   CreditScore          10000 non-null  int64
1   Geography            10000 non-null  object
2   Gender               10000 non-null  object
3   Age                 10000 non-null  int64
4   Tenure              10000 non-null  int64
5   Balance              10000 non-null  float64
6   NumOfProducts       10000 non-null  int64
7   HasCrCard           10000 non-null  int64
8   IsActiveMember      10000 non-null  int64
9   EstimatedSalary     10000 non-null  float64
10  Exited               10000 non-null  int64
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB
```

```
df["Geography"].unique()
array(['France', 'Spain', 'Germany'], dtype=object)
```

```
df["Gender"].unique()
array(['Female', 'Male'], dtype=object)
```

```
geo=pd.get_dummies(df["Geography"].drop_first=False)
```

```
geo.head()
France  Germany  Spain
0      1        0      0
1      0        0      1
2      1        0      0
3      1        0      0
4      0        0      1
```

```
gen=pd.get_dummies(df["Gender"].drop_first=False)
df=pd.concat([df, geo, gen], axis=1)
```

```
df
   CreditScore  Geography  Gender  Age  Tenure  Balance
NumOfProducts \
```


0	619	France	Female	42	2	0.00
1						
1	608	Spain	Female	41	1	83807.86
1						
2	502	France	Female	42	8	159660.80
3						
3	699	France	Female	39	1	0.00
2						
4	850	Spain	Female	43	2	125510.82
1						

...
...						
9995	771	France	Male	39	5	0.00
2						
9996	516	France	Male	35	10	57369.61
1						
9997	709	France	Female	36	7	0.00
1						
9998	772	Germany	Male	42	3	75075.31
2						
9999	792	France	Female	28	4	130142.79
1						

	HasCrCard	IsActiveMember	EstimatedSalary	Exited	France
--	-----------	----------------	-----------------	--------	--------

Germany \					
0	1	1	101348.88	1	1
0					
1	0	1	112542.58	0	0
0					
2	1	0	113931.57	1	1
0					
3	0	0	93826.63	0	1
0					
4	1	1	79084.10	0	0
0					

...
...					
9995	1	0	96270.64	0	1
0					
9996	1	1	101699.77	0	1
0					
9997	0	1	42085.58	1	1
0					
9998	1	0	92888.52	1	0
1					
9999	1	0	38190.78	0	1
0					

	Spain	Female	Male
0	0	1	0

```
1      1      1      0
2      0      1      0
3      0      1      0
4      1      1      0
...
9995    0      0      1
9996    0      0      1
9997    0      1      0
9998    0      0      1
9999    0      1      0
```

[10000 rows x 16 columns]

```
df.drop(["Geography","Gender"], axis=1, inplace=True)
```

```
df.head()
```

```
CreditScore  Age  Tenure  Balance  NumOfProducts  HasCrCard \
0           619   42      2    0.00              1          1
1           608   41      1  83807.86              1          0
2           502   42      8 159660.80              3          1
3           699   39      1    0.00              2          0
4           850   43      2 125510.82              1          1
```

```
IsActiveMember  EstimatedSalary  Exited  France  Germany  Spain
Female \
```

```
0           1    101348.88      1      1      0      0
1           1    112542.58      0      0      0      1
2           0    113931.57      1      1      0      0
3           0     93826.63      0      1      0      0
4           1     79084.10      0      0      0      1
```

```
Male
```

```
0      0
1      0
2      0
3      0
4      0
```

```
x=df.drop("Exited",axis=1)
```

x

```
CreditScore  Age  Tenure  Balance  NumOfProducts  HasCrCard \
0           619   42      2    0.00              1          1
1           608   41      1  83807.86              1          0
```

2	502	42	8	159660.80	3	1
3	699	39	1	0.00	2	0
4	850	43	2	125510.82	1	1
...
9995	771	39	5	0.00	2	1
9996	516	35	10	57369.61	1	1
9997	709	36	7	0.00	1	0
9998	772	42	3	75075.31	2	1
9999	792	28	4	130142.79	1	1

IsActive	Member	EstimatedSalary	France	Germany	Spain	Female
Male						
0	1	101348.88	1	0	0	1
0						
1	1	112542.58	0	0	1	1
0						
2	0	113931.57	1	0	0	1
0						
3	0	93826.63	1	0	0	1
0						
4	1	79084.10	0	0	1	1
0						

...
...						
9995	0	96270.64	1	0	0	0
1						
9996	1	101699.77	1	0	0	0
1						
9997	1	42085.58	1	0	0	1
0						
9998	0	92888.52	0	1	0	0
1						
9999	0	38190.78	1	0	0	1
0						

[[10000 rows x 13 columns]]

q=df[['Exited']]

q	
0	1
1	0
2	1
3	0
4	0

...	
9995	0
9996	0
9997	1
9998	1

```

9999 0
Name: Exited, Length: 10000, dtype: int64
df.shape
(10000, 14)
x.shape
(10000, 13)
y.shape
(10000,)
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=0)
x_train.shape
(8000, 13)
x_test.shape
(2000, 13)
y_test.shape
(2000,)
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_train
array([[ 0.16958176, -0.46460796,  0.00666099, ...,  1.74309049,
        1.09168714, -1.09168714],
       [-2.30455945,  0.30102557, -1.37744033, ..., -0.57369368,
        -0.91601335,  0.91601335],
       [-1.19119591, -0.94312892, -1.031415 , ..., -0.57369368,
        1.09168714, -1.09168714],
       ...,
       [ 0.9015152 , -0.36890377,  0.00666099, ..., -0.57369368,
        -0.91601335,  0.91601335],
       [-0.62420521, -0.08179119,  1.39076231, ...,  1.74309049,
        1.09168714, -1.09168714],
       [-0.28401079,  0.87525072, -1.37744033, ..., -0.57369368,
        1.09168714, -1.09168714]])
x_test = sc.transform(x_test)

```

x-test

```
array([[ -0.55204276, -0.36890377, 1.04473698, ..., -0.57369368,  
        1.09168714, -1.09168714],  
       [ -1.31490297,  0.10961719, -1.031415 , ..., -0.57369368,  
        1.09168714, -1.09168714],  
       [  0.57162971,  0.30102557, 1.04473698, ..., 1.74309049,  
        1.09168714, -1.09168714],  
       ...,  
       [ -0.74791227, -0.27319958, -1.37744033, ..., 1.74309049,  
        -0.91601335,  0.91601335],  
       [ -0.00566991, -0.46460796, -0.33936434, ..., -0.57369368,  
        -0.91601335,  0.91601335],  
       [ -0.79945688, -0.84742473, 1.04473698, ..., -0.57369368,  
        -0.91601335,  0.91601335]])
```